

CSE560 – Online Music Store Database System

Team - 04

Sathwick Kiran M S
UBID – sathwick

Kundan Anil Satkar
UBID - kundanan

I. Introduction

The music business has undergone a significant change in the digital age in the past two decades. As we are aware, online music retailers are now critical avenues for music distribution sales. To make decisions that facilitate growth, one must better understand sales behavior, customers, and the success of the business. This project will use SQL to analyze a data set of an online music store in a logical and systematic way. The primary purpose of the project is to provide insights that respond to critical business questions on sales, customer behavior, and how revenue is derived. Finding valuable insights is the full intent of this project. Ultimately, the project will hopefully make data-driven recommendations on improving customer experience satisfaction, retail process efficiency, and business profit.

Synopsis of Milestone 1

Our first milestone in the Online Music Store Database project involved building a foundation for our capabilities for managing and analysing data. We were designing and implementing a structured relational database management system, with a greater capacity for managing sites with larger datasets about music inventory, sales and user activity data. After determining the business problem i.e., the methods we were used to manage datasets like Excel were unacceptable in managing our growing online music store, we were able to justify the need for a database as opposed to traditional data management demonstrating key advantages with respect to data accuracy, scalability, analysis and security. We then identified the major user roles - customers, store managers and administrators - and documented their corresponding use cases to ensure the system would meet real-life requirements. Next we generated an E/R diagram and built eleven interrelated tables: Customer, Employee, Invoice, Invoice_Line, Track, Album, Artist, Media_Type, Genre, Playlist, and Playlist_Track. We built each table with primary keys and foreign keys to protect the integrity of the data and allow complex queries. In addition, we completed and tested some of our initial SQL queries in the following ways: querying the database for all artists, counting the total tracks by genre, rankings of total spending for all customers, determining the most 'productive' artist, and filtering playlists to only a specified number of tracks. These queries demonstrated that we could extract data from the database in a meaningful way. Overall, this phase was a good base for our milestone two implementation.

Problem Statement

If a new online music business fails to implement good management systems, tracking sales and customer information, and managing music inventories can seem

unmanageable, become susceptible to mistakes, and ultimately become impossible to grow. An organized system for diligence is necessary to ensure that you can track purchases, analyze customer preferences, and accurately maintain your sales records.

A structured management database system organizes these processes by managing large volumes of data, enabling complex analysis of the data, and ensuring integrity by establishing logical connections. The technology will also help a better business decision assistance, improve sales-process effectiveness, and enhance the customer experience by allowing and advancing automation, scalability, and precision.

II. Why there is a need for a database instead of Excel?

Efficient Data Management: A database efficiently organizes and deals with large quantities of information, such as customers, their musical preferences, and sales as the shop expands without limitation.

Data Accuracy and Consistency: A database assures that information is correct, organized, and consistent by enforcing rules such as primary and foreign keys.

Powerful Data Analysis: A database generates better analytics by utilizing SQL capabilities including facilities of SELECT, JOIN, GROUP BY, aggregated functions to facilitate decisions and enhance administration.

Optimized Storage and Maintenance: Data normalization removes redundancy and is more efficient for storage and updating by not needing to repeat the same data multiple times.

Enhanced Security and Recovery: Backup and recovery mechanisms both protect against data warehouse data loss and ensure that you can restore it quickly when you delete the wrong row or the system crashes.

III. Target Users

The database system supports three main types of users:

Customers:

- Browse music by category, artist, or genre.
- Place orders, check status, and manage their profiles and collections.

Store Managers:

- Add or update music albums and songs.
- Review customer orders and maintain inventory and product details.

Administrators:

- Monitor database health and ensure consistent, accurate data.
- Fix duplicated records and manage user accounts securely.

Real-Life Scenarios:

- **Business Owner:** Siddharth, the owner of an online music store (like iTunes), tracks sales trends, customer loyalty, and revenue patterns.
- **Data Analyst:** Sathwick analyzes sales data to find best-selling genres, artists, and peak sales periods, helping optimize promotions and stock.
- **Administrator:** Kundan manages database performance, optimizes slow queries, and ensures fast, reliable operations by maintaining clean, accurate data.

IV. Sample Database Design

The sample music store database was able to be created. This allows for the testing and debugging process to happen prior to working with a much larger data set. Below are examples of SQL statements demonstrating using all aspects of database management (creating, loading, altering, updating, and destroying).

1. Creating Sample table:

```
CREATE TABLE Track (
  TrackID INT PRIMARY KEY,
  Name VARCHAR (255),
  AlbumID INT,
  MediaTypeID INT,
  GenreID INT,
  Milliseconds INT,
  Bytes INT
);
```

2. Loading into the sample table:

```
INSERT INTO Track (TrackID, Name, AlbumID, MediaTypeID,
  GenreID, Milliseconds, Bytes)
VALUES (1, 'Sample Song', 10, 1, 5, 230000, 5120000);
```

3. Altering the Table:

```
ALTER TABLE Track ADD COLUMN UnitPrice NUMERIC(5,2);
```

4. Updating the Table:

```
UPDATE Track SET UnitPrice = 0.99
WHERE TrackID = 1;
```

5. Deleting the table:

```
DROP TABLE IF EXISTS Track;
```

V. Tables Description

| Table Name | Description |
|-------------------|--|
| Artist | Stores artist information including artist name and unique ID. |
| Album | Contains album details linked to artists. |
| Genre | Stores genre names and their identifiers. |
| Media_Type | Lists available media formats (e.g., MP3, AAC). |

| | |
|-----------------------|--|
| Track | Contains music track information including name, genre, album, and duration. |
| Playlist | Stores playlist details such as playlist name and ID. |
| Playlist_Track | Maps tracks to playlists using many-to-many relationships. |
| Customer | Stores customer details including name, contact info, and support rep ID. |
| Employee | Maintains employee information including job title, manager, and contact info. |
| Invoice | Records sales transactions including billing address and customer link. |
| Invoice_Line | Stores individual item sales from invoices including track, quantity, and price. |

Fig 1 – Base Table Description

VI. Foreign Key Action

| Table Name | Foreign Key | Referenced Table | On Delete Action |
|----------------|----------------|------------------|------------------|
| Album | artist_id | Artist | CASCADE |
| Track | album_id | Album | CASCADE |
| Track | genre_id | Genre | SET NULL |
| Track | media_type_id | Media_Type | SET NULL |
| Playlist_Track | playlist_id | Playlist | CASCADE |
| Playlist_Track | track_id | Track | CASCADE |
| Invoice | customer_id | Customer | CASCADE |
| Invoice | support_rep_id | Employee | SET NULL |
| Invoice_Line | invoice_id | Invoice | CASCADE |
| Invoice_Line | track_id | Track | CASCADE |
| Customer | support_rep_id | Employee | SET NULL |

Fig 2 – Foreign Keys Action

VII. Normalization Process & BCNF Validation

This tables explains the relation and their attributes before BCNF validation is done.

Now, when checking which tables violated BCNF and which tables satisfied this, we found that five tables that is.,

- Invoice
- Invoice_line
- Customer
- Employee
- Track

These tables were not in BCNF form, and we had to decompose this to satisfy this condition.

Decomposition:

1. Invoice-

Initial Schema –

```
CREATE TABLE public.invoice
(invoice_id character varying(30) NOT NULL,
```

customer_id character varying(30),
 invoice_date timestamp without time zone,
 billing_address character varying(120),
 billing_city character varying(30),
 billing_state character varying(30),
 billing_country character varying(30),
 billing_postal character varying(30),
 total double precision,
 PRIMARY KEY (invoice_id));

Violation-

- invoice_line_id → invoice_id, track_id, unit_price, quantity,
- track_id → unit_price which **violates BCNF** as track_id is not a superkey.

Decomposed -

```
CREATE TABLE customer_billing (
customer_id VARCHAR(30) PRIMARY KEY REFERENCES
customer(customer_id),
billing_address VARCHAR(120), billing_city VARCHAR(30),
billing_state VARCHAR(30), billing_country VARCHAR(30),
billing_postal VARCHAR(30));
INSERT INTO customer_billing
SELECT DISTINCT customer_id, billing_address, billing_city,
billing_state, billing_country, billing_postal
FROM invoice WHERE customer_id IS NOT NULL;
ALTER TABLE invoice DROP COLUMN billing_address,
DROP COLUMN billing_city,
DROP COLUMN billing_state, DROP COLUMN
billing_country, DROP COLUMN billing_postal;
```

2. Track-

Initial Schema-

```
CREATE TABLE public.track
(track_id character varying(30) NOT NULL,
name character varying(250),
album_id character varying(30),
media_type_id character varying(30),
genre_id character varying(30),
composer character varying(250),
milliseconds bigint,
bytes integer,
unit_price numeric,
PRIMARY KEY (track_id))
```

Violation-

- invoice_line_id → invoice_id, track_id, unit_price, quantity
- track_id → unit_price which **violates BCNF** as track_id is not a superkey.

Decomposed-

```
CREATE TABLE media_type_pricing (
media_type_id VARCHAR(30), unit_price NUMERIC,
PRIMARY KEY(media_type_id, unit_price));
INSERT INTO media_type_pricing
```

```
SELECT DISTINCT media_type_id, unit_price FROM track;
ALTER TABLE track DROP COLUMN unit_price;
```

3. Invoice Line-

Initial Schema-

```
CREATE TABLE public.invoice_line
(invoice_line_id character varying(30) NOT NULL,
invoice_id character varying(30),
track_id character varying(30),
unit_price numeric,
quantity integer,
PRIMARY KEY (invoice_line_id))
```

Violation-

- invoice_line_id → invoice_id, track_id, unit_price, quantity
- track_id → unit_price which **violates BCNF** as track_id is not a superkey.

Decomposition-

```
CREATE TABLE track_price (
track_id VARCHAR(30) PRIMARY KEY REFERENCES
track(track_id),
unit_price NUMERIC);
INSERT INTO track_price
SELECT DISTINCT track_id, unit_price FROM invoice_line;
ALTER TABLE invoice_line DROP COLUMN unit_price;
```

4. Employee-

Initial Schema-

```
CREATE TABLE public.employee
(
employee_id character varying(30) NOT NULL,
last_name character(50),
first_name character(50),
title character varying(250),
reports_to character varying(30),
levels character varying(10),
birth_date timestamp without time zone,
hire_date timestamp without time zone,
address character varying(120),
city character varying(50),
state character varying(50),
country character varying(30),
postal_code character varying(30),
phone character varying(30),
fax character varying(30),
email character varying(30),
PRIMARY KEY (employee_id))
```

Violation-

- track_id → all other attributes
- media_type_id → unit_price, genre_id → unit_price which **violates BCNF**.

Decomposition-

```
CREATE TABLE location (postal_code VARCHAR(30)
PRIMARY KEY, city VARCHAR(50), state VARCHAR(30),
country VARCHAR(30));
INSERT INTO location
SELECT DISTINCT postal_code, city, state, country FROM
employee WHERE postal_code IS NOT NULL;
ALTER TABLE employee DROP COLUMN city, DROP
COLUMN state, DROP COLUMN country;
```

5. Customer-

Initial Schema-

```
CREATE TABLE public.customer
(
customer_id character varying(30) NOT NULL,
first_name character(30),
last_name character(30),
company character varying(150),
address character varying(250),
city character varying(30),
state character varying(30),
country character varying(30),
postal_code character varying(30),
phone character varying(30),
fax character varying(30),
email character varying(30),
support_rep_id character varying(30),
PRIMARY KEY (customer_id));
```

Violation-

- customer_id → city, state, country
- postal_code → city, state, country which **violates BCNF** because postal_code is not a superkey in customer.

Decomposition-

```
CREATE TABLE customer_location (
postal_code VARCHAR(30) PRIMARY KEY, city
VARCHAR(30), state VARCHAR(30), country VARCHAR(30)
);
INSERT INTO customer_location
SELECT DISTINCT postal_code, city, state, country FROM
customer WHERE postal_code IS NOT NULL;
ALTER TABLE customer DROP COLUMN city, DROP
COLUMN state, DROP COLUMN country;
```

6. Album-

- Functional Dependency:
 - album_id → title, artist_id
- BCNF Status:
 - album_id is a superkey; no non-trivial FD violates BCNF.
- Justification-

The album table is already in BCNF because every album id uniquely determines its title and artist_id such that its not

possible to create a redundancy, partial dependency, and so no decomposition was needed.

7. Artist-

- Functional Dependency:
 - artist_id → name
- BCNF Status:
 - artist_id is a primary key directly determining the name.
- Justification

The artist table is in BCNF because the primary key artist_id can completely and uniquely determine the artist's name. No anomalies and no redundancies exist.

8. Genre-

- Functional Dependency:
 - genre_id → name
- BCNF Status:
 - Single simple key; no partial or transitive dependencies.
- Justification:

The genres table is entirely normalized because genre_id uniquely identifies the genre name and therefore does not contain evidence of redundancy, so this schema remains clean and simple.

9. Media Type-

- Functional Dependency:
 - media_type_id → name
- BCNF Status:
 - media_type_id is a primary key.
- Justification:

The media_types table is consistent with BCNF because each media_type_id directly corresponds to a single media type name, meaning there can be no duplication of data.

10. Playlist-

- Functional Dependency:
 - playlist_id → name
- BCNF Status:
 - playlist_id is a superkey; no issues with BCNF.
- Justification:

The playlist table is in normal form, because each playlist_id identifies a playlist name uniquely which prevents all redundancy and unnecessary data proliferation.

11. Playlist Track-

- Functional Dependency:
 - (playlist_id, track_id) → (composite PK for this join table)
- BCNF Status:
 - The combination (playlist_id, track_id) forms the full primary key; no partial or transitive dependencies exist.
- Justification:

The playlist_track table is already in BCNF, because since it represents a many-to-many relationship (between a playlist and a track), the composite primary key uniquely identifies each row.

After completing decomposition like mentioned above, the decomposition is completed, and all the tables are in BCNF.

| | | | |
|----------|-------------|--|-------------------|
| | | postal_code → city, state, country | Satisfies BCNF |
| location | postal_code | | |

Fig 3 – BCNF Analysis

VIII. Final BCNF Validated Schema

| Table Name | Keys and Attributes |
|---------------------------|---|
| album | PK: album_id album_id INT, title VARCHAR(150), artist_id INT (FK) |
| artist | PK: artist_id artist_id INT, name VARCHAR(100) |
| customer | PK: customer_id customer_id INT, first_name, last_name, email, phone, address, postal_code, support_rep_id INT (FK) |
| employee | PK: employee_id employee_id INT, first_name, last_name, title, reports_to, levels, birth_date, hire_date, address, postal_code, phone, fax, email |
| genre | PK: genre_id genre_id INT, name VARCHAR(50) |
| media_type | PK: media_type_id media_type_id INT, name VARCHAR(50) |
| track | PK: track_id track_id INT, name, album_id (FK), media_type_id (FK), genre_id (FK), composer, milliseconds, bytes |
| track_price | PK: track_id track_id INT (FK), unit_price NUMERIC |
| media_type_pricing | PK: media_type_id media_type_id INT (FK), unit_price NUMERIC |
| invoice | PK: invoice_id invoice_id INT, customer_id INT (FK), total NUMERIC, invoice_date DATE |
| customer_billing | PK: customer_id customer_id INT (FK), billing_address, billing_city, billing_state, billing_country, billing_postal_code |
| invoice_line | PK: invoice_line_id invoice_line_id INT, invoice_id INT (FK), track_id INT (FK), quantity INT |
| playlist | PK: playlist_id playlist_id INT, name VARCHAR(100) |
| playlist_track | PK: (playlist_id, track_id) playlist_id INT (FK), track_id INT (FK) |
| customer_location | PK: postal_code postal_code VARCHAR(30), city, state, country |
| location | PK: postal_code postal_code VARCHAR(30), city, state, country |

Fig 4 – Final Table after Decomposition

This confirms that all tables are in BCNF form. Each dependency has been explained.

IX. Finalized E/R Diagram

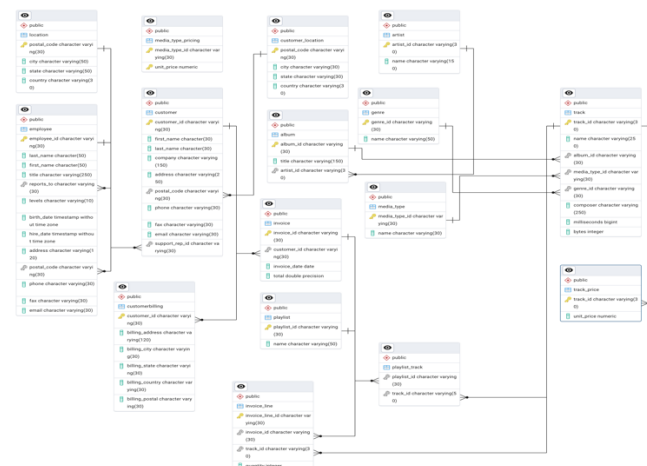


Fig 5 – Finalized E/R Diagram

| Relation | Candidate Key(s) | Functional Dependencies (FDs) | BCNF Status |
|--------------------|-------------------------|--|----------------------------|
| album | album_id | album_id → title, artist_id | Satisfies BCNF |
| artist | artist_id | artist_id → name | Satisfies BCNF |
| customer | customer_id | customer_id → name, email, support_rep_id, postal_code | Violates BCNF → Decomposed |
| employee | employee_id | employee_id → name, title, etc.; postal_code → city, state, country | Violates BCNF → Decomposed |
| genre | genre_id | genre_id → name | Satisfies BCNF |
| media_type | media_type_id | media_type_id → name | Satisfies BCNF |
| track | track_id | track_id → name, album_id, media_type_id, genre_id, etc.; media_type_id → unit_price | Violates BCNF → Decomposed |
| track_price | track_id | track_id → unit_price | Satisfies BCNF |
| invoice | invoice_id | invoice_id → customer_id, billing_info | Violates BCNF → Decomposed |
| invoice_line | invoice_line_id | invoice_line_id → invoice_id, track_id, quantity; track_id → unit_price | Violates BCNF → Decomposed |
| customer_billing | customer_id | customer_id → billing_address, billing_city, billing_state, billing_country, postal_code | Satisfies BCNF |
| media_type_pricing | media_type_id | media_type_id → unit_price | Satisfies BCNF |
| playlist | playlist_id | playlist_id → name | Satisfies BCNF |
| playlist_track | (playlist_id, track_id) | Composite key → determines each row | Satisfies BCNF |
| customer_location | postal_code | postal_code → city, state, country | Satisfies BCNF |

The final Entity-Relationship (E/R) diagram for our Online Music Store Database System consists of 16 normalized relations, derived from our previous selection of 11 separate tables. After applying the rules for BCNF normalization, 5 additional tables (and corresponding relationships) were added to eliminate transitive dependencies and ensure all determinants are a superkey.

Entity Design and Constraints

The database contains its basic entities: Customer, Employee, Invoice, Invoice_Line, Track, Album, Artist, Genre, Media_Type, Playlist, and Playlist_Track. Further relations were created to ensure BCNF compliance: Location, CustomerBilling, MediaTypePricing, and TrackPrice.

Key constraints include:

Primary Keys: Every table has a unique identifier (e.g. customer_id, track_id, invoice_id).

Foreign Keys: Enforces referential integrity (e.g. Invoice.customer_id → Customer, Track.track_id → TrackPrice, postal_code → Location).

Composite Keys: Used in Playlist_Track in the case of many-to-many relations.

NOT NULL: Applied to mandatory fields to encourage completeness.

BCNF and Handling of Redundancies

Functional dependencies, such as postal_code → city, state, country, have been broken up into Location.

The CustomerBilling relation separates the customer address information from the Invoice.

The pricing logic is normalized into TrackPrice and MediaTypePricing.

All relations in all relations are now BCNF compliant and level of redundancy avoided ensures consistent updating.

| Constraint Type | Purpose |
|----------------------------------|---|
| Primary Keys | Ensure uniqueness of each record (e.g., customer_id, track_id, invoice_id) |
| Foreign Keys | Enforce referential integrity between tables (e.g., invoice.customer_id → customer.customer_id) |
| CHECK Constraints | Validate domain rules (e.g., quantity > 0 in invoice_line, rating BETWEEN 1 AND 5) |
| Composite Keys | Represent many-to-many relations (e.g., playlist_track(playlist_id, track_id)) |
| Normalization Constraints | BCNF-compliant schema to avoid redundancy (e.g., TrackPrice, CustomerBilling, Location) |

Fig 6 -Constraints identified and implemented

X. Challenged Faced with Dataset and Solutions

We encountered; while working with larger datasets, we encountered notable performance bottlenecks, particularly during join and filter operations across multiple tables. One such challenge arose when executing the following query:

```
SELECT
  c.customer_id,
  c.first_name || ' ' || c.last_name AS customer_name,
  cl.city
FROM customer c
JOIN customer_location cl
  ON c.postal_code = cl.postal_code;
```

The join on postal_code between the customer and customer_location tables was significantly slower due to the absence of indexing on the join column. To address this, we created indexes on the postal_code attribute in both tables, which led to a substantial improvement in query speed.

In another query:

```
SELECT g.name, COUNT(t.track_id)
FROM genre g
JOIN track t ON g.genre_id = t.genre_id
GROUP BY g.name;
```

By indexing the genre_id and track_id columns, we enhanced the performance of aggregations and joins involving the track and genre tables. Additionally, we applied indexing to other frequently accessed fields like invoice_id and customer_id to boost overall query responsiveness. These indexing strategies were critical for maintaining performance and scalability as the dataset grew.

XI. Queries Executed

```
213 INSERT INTO customer_location (postal_code, city, state, country)
214 VALUES ('94105', 'San Francisco', 'CA', 'USA')
215 ON CONFLICT (postal_code) DO NOTHING;
216 ;
217 INSERT INTO customer (
218   customer_id,
219   first_name,
220   last_name,
221   postal_code,
222   support_rep_id
223 ) VALUES (
224   'CUST789',
225   'John',
226   'Doe',
227   '94105',
228   '3'
229 );
230 SELECT * FROM customer_location WHERE postal_code = '94105';
231
```




Fig 1 – Inserting into Table

This query first inserts a new location into the customer_location table, then adds a new customer linked to that postal code and finally retrieves the location details for postal code '94105'.

```

-- 3. DELETE (Remove one playlist-track mapping)
DELETE FROM playlist_track
WHERE playlist_id = '1'
AND track_id = '3402';

-- 4. DELETE (Remove one invoice line)
DELETE FROM invoice_line
WHERE invoice_line_id = '1';

-- 5. UPDATE (Increase price for a specific track)
UPDATE track_price
SET unit_price = unit_price * 1.10
WHERE track_id = '1';

```

Output Messages Notifications

FE 1

y returned successfully in 34 msec.

Fig 2 – Deleting from play_list

This query removes a specific mapping between a playlist and a track from the `playlist_track` table using `playlist_id = '1'` and `track_id = '3402'`. It deletes one row, indicating that this track is no longer part of the specified playlist.

```

-- 4. DELETE (Remove one invoice line)
DELETE FROM invoice_line
WHERE invoice_line_id = '1';

-- 5. UPDATE (Increase price for a specific
UPDATE track_price
SET unit_price = unit_price * 1.10
WHERE track_id = '1';

-- 6. UPDATE (Adjust quantity on an invoice)
UPDATE invoice_line
SET quantity = 2

```

Output Messages Notifications

TE 1

y returned successfully in 42 msec.

Fig 3 – Deleting from invoice_line

This query deletes '1' from `invoice_line`.

```

432
433 -- 5. UPDATE Increase price for a specific track
434 UPDATE track_price
435 SET unit_price = unit_price * 1.10
436 WHERE track_id = '1';
437 select * from track_price;
438

```

Data Output Messages Explain X Notifications

UPDATE 1

Query returned successfully in 58 msec.

Fig 4 – Updating price

This SQL UPDATE statement increases the `unit_price` of the track with `track_id = '1'` by 10%. It multiplies the current price by 1.10 and updates the `track_price` table accordingly. The operation completed successfully, affecting one row.

```

239 -- 7. SELECT with JOIN Customers and their cities
240 SELECT
241 c.customer_id,
242 c.first_name || ' ' || c.last_name AS customer_name,
243 c.city
244 FROM customer c
245 JOIN customer_location cl
246 ON c.postal_code = cl.postal_code;
247

```

Data Output Messages Notifications

| | customer_id character varying (30) | customer_name text | city character varying (30) |
|----|---------------------------------------|-----------------------|--------------------------------|
| 1 | customer_id | first_name last_name | city |
| 2 | 1 | Luis Gonçalves | São José dos Campos |
| 3 | 2 | Leonie Köhler | Stuttgart |
| 4 | 3 | François Tremblay | Montréal |
| 5 | 4 | Bjørn Hansen | Oslo |
| 6 | 5 | František Wichterlová | Prague |
| 7 | 6 | Helena Holý | Prague |
| 8 | 7 | Astrid Gruber | Vienne |
| 9 | 8 | Daan Peeters | Brussels |
| 10 | 9 | Kara Nielsen | Copenhagen |
| 11 | 10 | Eduardo Martins | São Paulo |
| 12 | 11 | Alexandre Rocha | São Paulo |
| 13 | 12 | Roberto Almeida | Rio de Janeiro |

Fig 5 – Joining customer and city

This query joins the `customer` and `customer_location` tables using `postal_code` to display each customer's name along with their city. It demonstrates how relational joins help in enriching customer data with geographic context.

```

247 -- 8. SELECT with GROUP BY Number of tracks per genre
248 SELECT
249 g.name AS genre_name,
250 COUNT(t.track_id) AS track_count
251 FROM genre g
252 JOIN track t
253 ON g.genre_id = t.genre_id
254 GROUP BY g.name;
255

```

Data Output Messages Notifications

| | genre_name character varying (50) | track_count bigint |
|----|--------------------------------------|-----------------------|
| 1 | Heavy Metal | 28 |
| 2 | TV Shows | 93 |
| 3 | Latin | 579 |
| 4 | Electronica/Dance | 30 |
| 5 | R&B/Soul | 61 |
| 6 | Opera | 1 |
| 7 | Comedy | 17 |
| 8 | Classical | 74 |
| 9 | Pop | 48 |
| 10 | Easy Listening | 24 |
| 11 | Alternative | 40 |
| 12 | Drama | 64 |
| 13 | Sci Fi & Fantasy | 26 |
| 14 | Rock | 1297 |
| 15 | World | 28 |
| 16 | Reggae | 58 |
| 17 | Blues | 81 |

Fig 6 – Selecting with group by

This query retrieves the number of tracks available in each genre by grouping and counting track IDs. It helps identify which genres have the most music content in the database.

```

256
257 -- 9. SELECT with ORDER BY Top 5 longest tracks
258 SELECT
259 track_id,
260 name,
261 milliseconds
262 FROM track
263 ORDER BY milliseconds DESC
264 LIMIT 5;
265

```

Data Output Messages Notifications

| | track_id [PK] character varying (30) | name character varying (250) | milliseconds bigint |
|---|---|---------------------------------|------------------------|
| 1 | 2820 | Occupation / Precipice | 5286953 |
| 2 | 3224 | Through a Looking Glass | 5088838 |
| 3 | 3244 | Greetings from Earth, Pt... | 2960293 |
| 4 | 3242 | The Man With Nine Lives | 2956998 |
| 5 | 3227 | Battlestar Galactica, Pt. 2 | 2956081 |

Fig 7 – Select with Order By

This query returns the top 5 longest tracks by ordering all tracks in descending order of duration (milliseconds). It is useful for highlighting extended compositions or soundtracks in the database.

```

481 -- Query: Total revenue spent per customer per genre
482
483 SELECT
484     c.customer_id,
485     g.name AS genre_name,
486     SUM(il.quantity * tp.unit_price) AS total_spent
487 FROM customer c
488 JOIN invoice i ON c.customer_id = i.customer_id
489 JOIN invoice_line il ON i.invoice_id = il.invoice_id
490 JOIN track t ON il.track_id = t.track_id
491 JOIN track_price tp ON t.track_id = tp.track_id
492 JOIN genre g ON t.genre_id = g.genre_id
493 GROUP BY c.customer_id, g.name
494 ORDER BY total_spent DESC;
495

```

| | customer_id | genre_name | total_spent |
|----|-------------|------------|-------------|
| 1 | 6 | Rock | 75.24 |
| 2 | 3 | Rock | 74.25 |
| 3 | 46 | Rock | 71.28 |
| 4 | 1 | Rock | 71.28 |
| 5 | 34 | Rock | 67.32 |
| 6 | 5 | Rock | 66.33 |
| 7 | 37 | Rock | 64.35 |
| 8 | 57 | Rock | 60.39 |
| 9 | 51 | Rock | 59.40 |
| 10 | 21 | Rock | 58.617900 |
| 11 | 27 | Rock | 58.41 |
| 12 | 53 | Rock | 56.637900 |
| 13 | 64 | Rock | 56.637900 |

Fig 8 – Total revenue spent/customer

This query calculates the total amount each customer spent on each genre by joining invoices, invoice lines, tracks, and pricing details.

It groups results by customer and genre, showing the highest-spending customers per genre.

```

495
496 -- 12. Most Diverse Customers based on Number of Genres Purchased
497 SELECT
498     c.customer_id,
499     c.first_name || ' ' || c.last_name AS customer_name,
500     COUNT(DISTINCT t.genre_id) AS genres_purchased
501 FROM customer c
502 JOIN invoice i ON c.customer_id = i.customer_id
503 JOIN invoice_line il ON i.invoice_id = il.invoice_id
504 JOIN track t ON il.track_id = t.track_id
505 GROUP BY c.customer_id, customer_name
506 ORDER BY genres_purchased DESC
507 LIMIT 5;
508

```

| | customer_id | customer_name | genres_purchased |
|---|-------------|------------------|------------------|
| 1 | 2 | Leonie Köhler | 14 |
| 2 | 35 | Madalena Samp... | 13 |
| 3 | 30 | Edward Francis | 13 |
| 4 | 22 | Heather Leacock | 13 |
| 5 | 44 | Terhi Hämäläinen | 13 |

Fig 9 – Most Diverse Customers

This query identifies the most diverse customers by counting how many unique genres each has purchased from. It joins customer, invoice, invoice_line, and track tables to compute genre diversity per user. The result ranks customers based on the count of distinct genres, showing the top 5.

```

678 SELECT
679     c.customer_id,
680     c.first_name || ' ' || c.last_name AS customer_name,
681     SUM(il.quantity * tp.unit_price) AS total_spent
682 FROM customer c
683 JOIN invoice i ON c.customer_id = i.customer_id
684 JOIN invoice_line il ON i.invoice_id = il.invoice_id
685 JOIN track t ON il.track_id = t.track_id
686 JOIN track_price tp ON t.track_id = tp.track_id
687 GROUP BY c.customer_id, customer_name
688 HAVING SUM(il.quantity * tp.unit_price) > (
689     SELECT AVG(total)
690     FROM (
691         SELECT SUM(il.quantity * tp.unit_price) AS total
692         FROM invoice i
693         JOIN invoice_line il ON i.invoice_id = il.invoice_id
694         JOIN track t ON il.track_id = t.track_id
695         JOIN track_price tp ON t.track_id = tp.track_id
696         GROUP BY i.customer_id
697     ) AS sub
698 )
699 ORDER BY total_spent DESC
700 LIMIT 5;
701

```

| | customer_id | customer_name | total_spent |
|---|-------------|-----------------------|-------------|
| 1 | 5 | František Wichterlová | 144.54 |
| 2 | 6 | Helena Holý | 128.70 |
| 3 | 46 | Hugh O'Reilly | 114.84 |
| 4 | 58 | Manoj Pareek | 111.87 |
| 5 | 1 | Luís Gonçalves | 108.90 |

Fig 10 – SubQuery

This query retrieves the top 5 customers whose total spending is above the average spending of all customers.

XII. Query Execution Analysis

1. Joining Customer, invoice and invoice_line table-

EXPLAIN ANALYZE

```

SELECT c.customer_id, c.first_name, c.last_name,
SUM(il.quantity)
FROM customer c
JOIN invoice i ON c.customer_id = i.customer_id
JOIN invoice_line il ON i.invoice_id = il.invoice_id
GROUP BY c.customer_id, c.first_name, c.last_name
ORDER BY SUM(il.quantity) DESC;

```

```

704 EXPLAIN ANALYZE
705 SELECT g.name AS genre_name, COUNT(t.track_id) AS num_tracks
706 FROM track t
707 JOIN genre g ON t.genre_id = g.genre_id
708 JOIN album a ON t.album_id = a.album_id
709 GROUP BY g.name
710 ORDER BY num_tracks DESC;
711
712

```

| | QUERY PLAN |
|----|---|
| 1 | Sort (cost=131.84..131.91 rows=27 width=17) (actual time=7.293..7.299 rows=25 loops=1) |
| 2 | Sort Key: (count(t.track_id)) DESC |
| 3 | Sort Method: quicksort Memory: 25kB |
| 4 | → HashAggregate (cost=130.93..131.20 rows=27 width=17) (actual time=7.261..7.271 rows=25 loops=1) |
| 5 | Group Key: g.name |
| 6 | Batches: 1 Memory Usage: 24kB |
| 7 | → Hash Join (cost=12.44..113.42 rows=3503 width=13) (actual time=0.559..5.672 rows=3503 loops=1) |
| 8 | Hash Cond: (i.album_id = a.album_id) |
| 9 | → Hash Join (cost=1.61..93.28 rows=3503 width=16) (actual time=0.061..3.294 rows=3503 loops=1) |
| 10 | Hash Cond: (i.genre_id = g.genre_id) |
| 11 | → Seq Scan on track t (cost=0.00..81.03 rows=3503 width=9) (actual time=0.011..1.402 rows=3503 loops=1) |
| 12 | → Hash (cost=1.27..1.27 rows=27 width=11) (actual time=0.040..0.041 rows=27 loops=1) |
| 13 | Buckets: 1024 Batches: 1 Memory Usage: 10kB |
| 14 | → Seq Scan on genre g (cost=0.00..1.27 rows=27 width=11) (actual time=0.016..0.019 rows=27 loops=1) |
| 15 | → Hash (cost=5.48..5.48 rows=348 width=3) (actual time=0.487..0.488 rows=348 loops=1) |
| 16 | Buckets: 1024 Batches: 1 Memory Usage: 21kB |
| 17 | → Seq Scan on album a (cost=0.00..5.48 rows=348 width=3) (actual time=0.005..0.402 rows=348 loops=1) |
| 18 | Planning Time: 0.757 ms |
| 19 | Execution Time: 7.377 ms |

Fig 1 – Before Indexing

CREATE INDEX idx_invoice_customer_id ON invoice(customer_id);

CREATE INDEX idx_invoice_line_invoice_id ON invoice_line(invoice_id);


```

704 --
705 --
706 -- EXPLAIN ANALYZE
707 -- SELECT g.name AS genre_name, COUNT(t.track_id) AS num_tracks
708 -- FROM track t
709 -- JOIN genre g ON t.genre_id = g.genre_id
710 -- JOIN album a ON t.album_id = a.album_id
711 -- GROUP BY g.name
712 -- ORDER BY num_tracks DESC;

```

Query Plan

```

1  Sort (cost=131.84..131.91 rows=27 width=17) (actual time=4.325..4.334 rows=25 loops=1)
2   Sort Key: (count(t.track_id)) DESC
3   Sort Method: quicksort Memory: 25kB
4   => HashAggregate (cost=130.93..131.20 rows=27 width=17) (actual time=4.286..4.300 rows=25 loops=1)
5     Group Key: g.name
6     Batches: 1 Memory Usage: 24kB
7     => Hash Join (cost=12.44..113.42 rows=3503 width=13) (actual time=0.143..3.303 rows=3503 loops=1)
8       Hash Cond: ((t.album_id)::text = (a.album_id)::text)
9       => Hash Join (cost=1.61..93.28 rows=3503 width=16) (actual time=0.062..2.076 rows=3503 loops=1)
10        Hash Cond: ((t.genre_id)::text = (g.genre_id)::text)
11        => Seq Scan on track t (cost=0.00..81.03 rows=3503 width=9) (actual time=0.012..0.905 rows=3503 loops=1)
12        => Hash (cost=1.27..1.27 rows=27 width=11) (actual time=0.021..0.022 rows=27 loops=1)
13          Buckets: 1024 Batches: 1 Memory Usage: 10kB
14          => Seq Scan on genre g (cost=0.00..1.27 rows=27 width=11) (actual time=0.006..0.009 rows=27 loops=1)
15        => Hash (cost=6.48..6.48 rows=348 width=3) (actual time=0.074..0.075 rows=348 loops=1)
16          Buckets: 1024 Batches: 1 Memory Usage: 21kB
17          => Seq Scan on album a (cost=0.00..6.48 rows=348 width=3) (actual time=0.004..0.033 rows=348 loops=1)
18 Planning Time: 0.336 ms
19 Execution Time: 4.422 ms

```

Fig 2 – After Indexing

2. Joining track, genre and album with join and group by

```

EXPLAIN ANALYZE
SELECT g.name AS genre_name, COUNT(t.track_id) AS
num_tracks
FROM track t
JOIN genre g ON t.genre_id = g.genre_id
JOIN album a ON t.album_id = a.album_id
GROUP BY g.name
ORDER BY num_tracks DESC;

```

```

724 --
725 --
726 -- SELECT g.name AS genre_name, COUNT(t.track_id) AS num_tracks
727 -- FROM track t
728 -- JOIN genre g ON t.genre_id = g.genre_id
729 -- JOIN album a ON t.album_id = a.album_id
730 -- GROUP BY g.name
731 -- ORDER BY num_tracks DESC;

```

Query Plan

```

1  Sort (cost=131.84..131.91 rows=27 width=17) (actual time=2.913..2.917 rows=25 loops=1)
2   Sort Key: (count(t.track_id)) DESC
3   Sort Method: quicksort Memory: 25kB
4   => HashAggregate (cost=130.93..131.20 rows=27 width=17) (actual time=2.897..2.902 rows=25 loops=1)
5     Group Key: g.name
6     Batches: 1 Memory Usage: 24kB
7     => Hash Join (cost=12.44..113.42 rows=3503 width=13) (actual time=0.146..2.255 rows=3503 loops=1)
8       Hash Cond: ((t.album_id)::text = (a.album_id)::text)
9       => Hash Join (cost=1.61..93.28 rows=3503 width=16) (actual time=0.042..1.373 rows=3503 loops=1)
10        Hash Cond: ((t.genre_id)::text = (g.genre_id)::text)
11        => Seq Scan on track t (cost=0.00..81.03 rows=3503 width=9) (actual time=0.011..0.600 rows=3503 loops=1)
12        => Hash (cost=1.27..1.27 rows=27 width=11) (actual time=0.022..0.023 rows=27 loops=1)
13          Buckets: 1024 Batches: 1 Memory Usage: 10kB
14          => Seq Scan on genre g (cost=0.00..1.27 rows=27 width=11) (actual time=0.007..0.010 rows=27 loops=1)
15        => Hash (cost=6.48..6.48 rows=348 width=3) (actual time=0.093..0.093 rows=348 loops=1)
16          Buckets: 1024 Batches: 1 Memory Usage: 21kB
17          => Seq Scan on album a (cost=0.00..6.48 rows=348 width=3) (actual time=0.003..0.042 rows=348 loops=1)
18 Planning Time: 0.540 ms
19 Execution Time: 2.979 ms

```

Fig 3 - Before Indexing

```

CREATE INDEX idx_track_genre_id ON track(genre_id);
CREATE INDEX idx_track_album_id ON track(album_id);

```

```

731 --
732 --
733 -- CREATE INDEX idx_track_genre_id ON track(genre_id);
734 -- CREATE INDEX idx_track_album_id ON track(album_id);

```

Query Plan

```

1  Sort (cost=131.84..131.91 rows=27 width=17) (actual time=3.014..3.017 rows=25 loops=1)
2   Sort Key: (count(t.track_id)) DESC
3   Sort Method: quicksort Memory: 25kB
4   => HashAggregate (cost=130.93..131.20 rows=27 width=17) (actual time=2.996..3.001 rows=25 loops=1)
5     Group Key: g.name
6     Batches: 1 Memory Usage: 24kB
7     => Hash Join (cost=12.44..113.42 rows=3503 width=13) (actual time=0.204..2.303 rows=3503 loops=1)
8       Hash Cond: ((t.album_id)::text = (a.album_id)::text)
9       => Hash Join (cost=1.61..93.28 rows=3503 width=16) (actual time=0.043..1.325 rows=3503 loops=1)
10        Hash Cond: ((t.genre_id)::text = (g.genre_id)::text)
11        => Seq Scan on track t (cost=0.00..81.03 rows=3503 width=9) (actual time=0.009..0.512 rows=3503 loops=1)
12        => Hash (cost=1.27..1.27 rows=27 width=11) (actual time=0.020..0.020 rows=27 loops=1)
13          Buckets: 1024 Batches: 1 Memory Usage: 10kB
14          => Seq Scan on genre g (cost=0.00..1.27 rows=27 width=11) (actual time=0.006..0.009 rows=27 loops=1)
15        => Hash (cost=6.48..6.48 rows=348 width=3) (actual time=0.155..0.155 rows=348 loops=1)
16          Buckets: 1024 Batches: 1 Memory Usage: 21kB
17          => Seq Scan on album a (cost=0.00..6.48 rows=348 width=3) (actual time=0.004..0.101 rows=348 loops=1)
18 Planning Time: 0.794 ms
19 Execution Time: 3.083 ms

```

Fig 4 – After Indexing

3. Playlist revenue Ranking

```

EXPLAIN ANALYZE
SELECT
p.playlist_id,
p.name AS playlist_name,
COUNT(pt.track_id) AS num_tracks,
SUM(il.quantity * tp.unit_price) AS total_revenue,
RANK() OVER (ORDER BY SUM(il.quantity * tp.unit_price)
DESC) AS revenue_rank
FROM playlist p
JOIN playlist_track pt ON p.playlist_id = pt.playlist_id
JOIN invoice_line il ON pt.track_id = il.track_id
JOIN track_price tp ON il.track_id = tp.track_id
GROUP BY p.playlist_id, p.name
HAVING SUM(il.quantity * tp.unit_price) > 1000
ORDER BY total_revenue DESC;

```

COUNT(pt.track_id) AS num_tracks,
SUM(il.quantity * tp.unit_price) AS total_revenue,
RANK() OVER (ORDER BY SUM(il.quantity * tp.unit_price)
DESC)

AS revenue_rank

FROM playlist p
JOIN playlist_track pt ON p.playlist_id = pt.playlist_id
JOIN invoice_line il ON pt.track_id = il.track_id
JOIN track_price tp ON il.track_id = tp.track_id
GROUP BY p.playlist_id, p.name
HAVING SUM(il.quantity * tp.unit_price) > 1000
ORDER BY total_revenue DESC;

```

1  WindowAgg (cost=518.46..518.57 rows=6 width=63) (actual time=20.178..20.189 rows=3 loops=1)
2   => Sort (cost=518.46..518.48 rows=6 width=63) (actual time=20.139..20.144 rows=3 loops=1)
3   Sort Key: (sum((il.quantity)::numeric * tp.unit_price)) DESC
4   Sort Method: quicksort Memory: 25kB
5   => HashAggregate (cost=518.10..518.39 rows=6 width=55) (actual time=20.123..20.133 rows=3 loops=1)
6     Group Key: p.playlist_id
7     Filter: (sum((il.quantity)::numeric * tp.unit_price)) > '1000'::numeric
8     Batches: 1 Memory Usage: 24kB
9     Rows Removed by Filter: 10
10    => Hash Join (cost=272.68..441.91 rows=6095 width=28) (actual time=9.598..14.093 rows=11558 loops=1)
11      Hash Cond: ((t.track_id)::text = (pt.track_id)::text)
12      => Seq Scan on invoice_line il (cost=0.00..78.56 rows=4756 width=8) (actual time=0.016..0.929 rows=4756 loops=1)
13      => Hash (cost=216.83..216.53 rows=4492 width=28) (actual time=9.535..9.539 rows=4476 loops=1)
14        Buckets: 8192 Batches: 1 Memory Usage: 308kB
15        => Hash Join (cost=92.06..216.53 rows=4492 width=28) (actual time=1.061..7.653 rows=4476 loops=1)
16          Hash Cond: ((pt.playlist_id)::text = (p.playlist_id)::text)
17          => Hash Join (cost=50.63..200.68 rows=4492 width=15) (actual time=1.040..5.653 rows=4476 loops=1)
18            Hash Cond: ((pt.track_id)::text = (tp.track_id)::text)
19            => Seq Scan on playlist_track pt (cost=0.00..127.15 rows=8715 width=6) (actual time=0.008..0.008 rows=8715 loops=1)
20            => Hash (cost=28.06..28.06 rows=1806 width=9) (actual time=0.881..0.892 rows=1806 loops=1)
21              Buckets: 2048 Batches: 1 Memory Usage: 90kB
22              => Seq Scan on track_price tp (cost=0.00..28.06 rows=1806 width=9) (actual time=0.007..0.007 rows=1806 loops=1)
23            => Hash (cost=1.19..1.19 rows=19 width=15) (actual time=0.015..0.016 rows=19 loops=1)
24              Buckets: 1024 Batches: 1 Memory Usage: 9kB
25              => Seq Scan on playlist p (cost=0.00..1.19 rows=19 width=15) (actual time=0.005..0.008 rows=19 loops=1)
26 Planning Time: 2.189 ms
27 Execution Time: 20.361 ms

```

Fig 5 – Before Indexing

```

CREATE INDEX idx_playlist_track_playlist_id
ON playlist_track(playlist_id);
CREATE INDEX idx_invoice_line_track_id
ON invoice_line(track_id);
CREATE INDEX idx_track_price_track_id
ON track_price(track_id);

```

```

747 --
748 -- JOIN invoice_line il ON pt.track_id = il.track_id
749 -- JOIN track_price tp ON pt.track_id = tp.track_id
750 -- HAVING SUM(il.quantity * tp.unit_price) > 1000
751 -- ORDER BY total_revenue DESC;

```

Query Plan

```

5  => HashAggregate (cost=518.10..518.39 rows=6 width=55) (actual time=19.428..19.435 rows=3 loops=1)
6     Group Key: p.playlist_id
7     Filter: (sum((il.quantity)::numeric * tp.unit_price)) > '1000'::numeric
8     Batches: 1 Memory Usage: 24kB
9     Rows Removed by Filter: 10
10    => Hash Join (cost=272.68..441.91 rows=6095 width=28) (actual time=9.929..14.157 rows=11558 loops=1)
11      Hash Cond: ((t.track_id)::text = (pt.track_id)::text)
12      => Seq Scan on invoice_line il (cost=0.00..78.56 rows=4756 width=8) (actual time=0.011..0.868 rows=4756 loops=1)
13      => Hash (cost=216.83..216.53 rows=4492 width=28) (actual time=9.883..9.887 rows=4476 loops=1)
14        Buckets: 8192 Batches: 1 Memory Usage: 308kB
15        => Hash Join (cost=92.06..216.53 rows=4492 width=28) (actual time=0.390..7.891 rows=4476 loops=1)
16          Hash Cond: ((pt.playlist_id)::text = (p.playlist_id)::text)
17          => Hash Join (cost=50.63..200.68 rows=4492 width=15) (actual time=0.369..5.798 rows=4476 loops=1)
18            Hash Cond: ((pt.track_id)::text = (tp.track_id)::text)
19            => Seq Scan on playlist_track pt (cost=0.00..127.15 rows=8715 width=6) (actual time=0.005..0.005 rows=8715 loops=1)
20            => Hash (cost=28.06..28.06 rows=1806 width=9) (actual time=0.358..0.359 rows=1806 loops=1)
21              Buckets: 2048 Batches: 1 Memory Usage: 90kB
22              => Seq Scan on track_price tp (cost=0.00..28.06 rows=1806 width=9) (actual time=0.008..0.008 rows=1806 loops=1)
23            => Hash (cost=1.19..1.19 rows=19 width=15) (actual time=0.009..0.009 rows=19 loops=1)
24              Buckets: 1024 Batches: 1 Memory Usage: 9kB
25              => Seq Scan on playlist p (cost=0.00..1.19 rows=19 width=15) (actual time=0.003..0.004 rows=19 loops=1)
26 Planning Time: 4.010 ms
27 Execution Time: 19.588 ms

```

Fig 6 – After Indexing